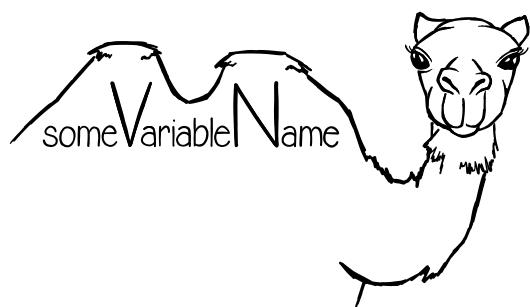


or

or

or

...



000

Kleene's Star is not Blah Blah

Computers are great at repetitive tasks. The most beautiful and concise way of expressing repetition is Kleene's star, namely the symbol * which stands for **zero or more repetitions**. In Computer Science theory, the letter A followed by the Kleene Star, A*, either means

- "" (the word of length 0), or
- "A" (the word consisting of one A), or
- "AA" (the word consisting of exactly two As), etc

A related but different concept is a *wildcard* where the star is a placeholder (instead of a sign of repetition). For example *.txt selects all files that end in "txt". Kleene's star is not common in daily life, but wildcards are: A prominent example is "blah blah".

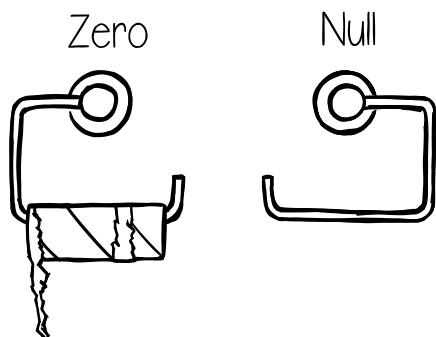
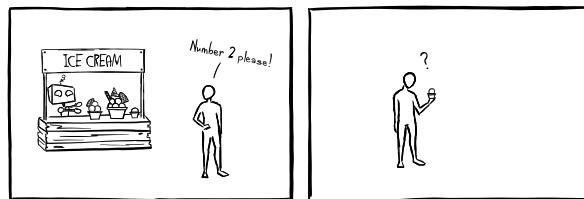
002

CamelCase, or why coding conventions matter

One problem that often arises when writing code is the naming of variables. Often multi-word names are used for variables but in most programming languages it is not possible to use spaces in a variable name.

One programming style that tries to structure this is called CamelCase. It is the practice of writing each word in the middle of a phrase with a Capital letter. This greatly improves readability.

thisGreatlyImprovesReadabilityComparedTo
whenewouldusenocapitalizationatall



005

**"Should array indices start at 0 or 1?
My compromise of 0.5 was rejected without,
I thought, proper consideration."**

An array is a data structure to store some elements with an index for each. This coins the question: 'what is the smallest index?'

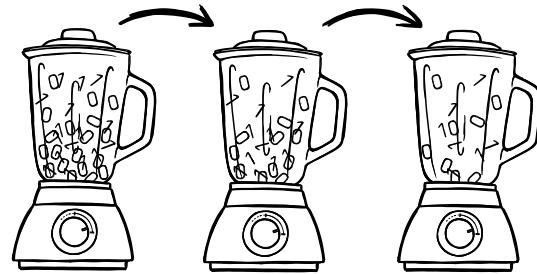
Most programming languages start an array at index 0 because it is the smallest natural number that can be represented in a binary system. But there are some rare outliers that start with index 1.

If you want to figure out whether the person in front of you is a computer scientist or not, there is one easy way to find it out: Show a bunch of things and let them count them. The true computer scientist will start to count with a zero.

007

Zero is not null

The null value has a special meaning in computer science. It is used to declare that a value is not yet defined or known. The specialty of null is that although it shows us that a value is missing, it itself is a value. So, it is possible to compare two nonexistent values. Suppose we have data about a patient; it does make a difference whether the patient has no disease (0) or is not yet diagnosed (null).



Conditions		OR	XOR
A	B		
○	○	○	○
●	○	●	●
○	●	●	●
●	●	●	○

Mincing Bytes (or the Art of Hashing)

010

By thoroughly mincing and blending a file's content several times, each time keeping only a small part of it, one can extract a bit pattern called a hash. For all practical purposes, such a hash value is unique to the underlying data.

One can not overstate the importance of hash computations for cryptographic algorithms as well as for efficient data management.

In cryptography, hashing yields a hard to forge "fingerprint" so that an attacker cannot change the terms of an electronic contract without also changing the fingerprint value, for example.

In data management, hashing serves to "spread out" items in space, like assigning a unique slot number to arbitrary data items, with a low chance of assigning that number twice.

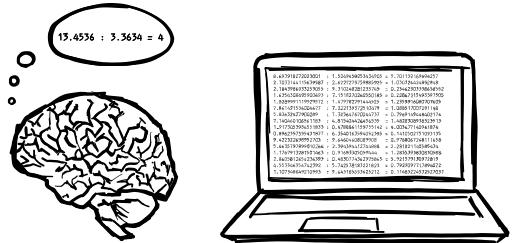
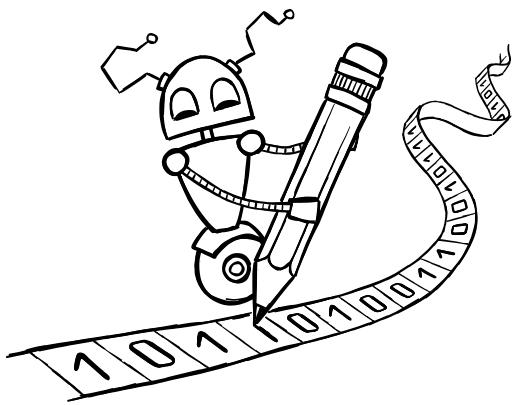
There are two different kinds of OR

013

The question whether an "or" can also be interpreted as an "and" is a common source of confusion. While as humans we can often resolve this confusion by looking at the context, computers require a strict definition.

In mathematical logic, the foundation of Computer Science, it is therefore distinguished between "OR" and "exclusive OR". The latter is often abbreviated as XOR. In English, we can express an XOR by using an "either ... or ..." construction.

The OR operator is always defined as an "and/or". It is fulfilled if at least one of its conditions is fulfilled. It is therefore also called "inclusive or". An exclusive or on the other hand is fulfilled if (and only if) exactly one of the conditions is fulfilled. If both of its conditions are met, the exclusive or is no longer fulfilled.



The Essence of Computation: The Turing Machine

019

The Turing machine is an imaginary model of a machine. It describes what can be computed by following an algorithm. It prints information in a coded form on a tape. This tape is divided into squares. Each of these squares contains either a 0 or a 1.

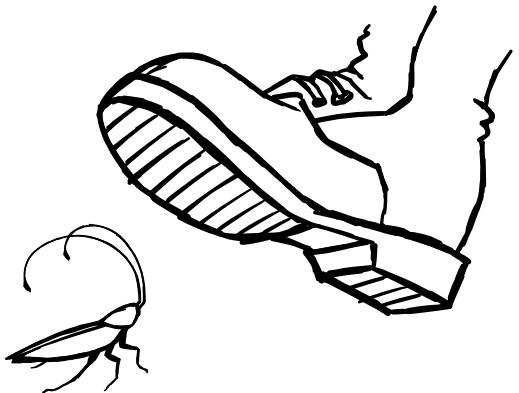
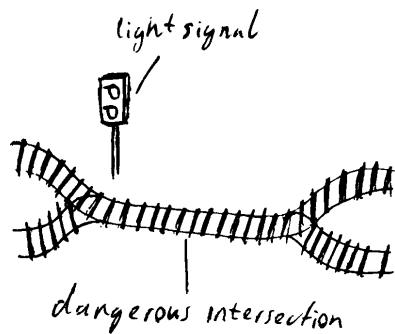
As a state machine, the Turing machine reads the content of the currently selected square and executes the algorithm associated with its current state. Every possible state of the Turing machine has an algorithm associated. Depending on the algorithm, the machine modifies the content of a square and moves the tape to the left or the right.

The Turing machine is a quite simple concept but it turns out, it contains the essence of computation: Every problem that can be computed, can theoretically also be computed by a Turing machine. This makes the Turing machine an important concept in theoretical computer science.

The more FLOPS the better

023

FLOPS is a performance measurement indicating how many floating point operations can be calculated in a second. It is used to measure the performance of computers and their processors. A floating point operation describes a mathematical procedure using decimal numbers. Despite the human brain being immensely powerful its capabilities when doing decimal calculations is limited. It is estimated at around 0.01 FLOPS, roughly 1 calculation every 100 seconds. In comparison, a modern smartphone has about 10^{12} FLOPS (1 Teraflop), while the fastest super computer has around 100×10^{15} FLOPS (100 Petaflops). Historically, the power of processors has been increasing exponentially (according to Moore's Law).



Atomic Operations

024

Imagine a railway intersection where only one train can pass at a time. If any other train tries to pass the intersection while it's already being used, a crash happens.

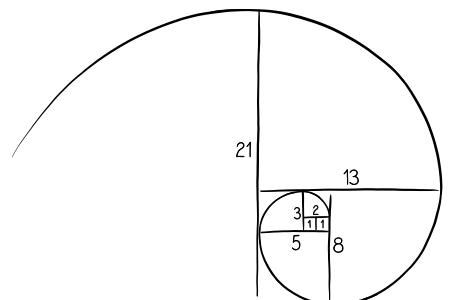
The same thing can happen in a computer, where we might think that all things happen ordered and sequentially. But in a world where everything is parallelized this can not be guaranteed and as soon as a file gets changed and read at the same time, things can go wrong.

Similar to a light signal for the railway intersection which would only allow one train to pass at a time, computers have something called *atomic operations*. As the name suggests, these are the smallest possible operations and therefore always done in their entirety with no space for anything in between, thus preventing crashes.

We all hate Bugs

025

Bugs creep into your house without notice, they eat your food and clog your sink. They hide in corners and behind large furniture and are most of the times hard to find. If you find a bug that was bothering you for a long time the feeling of success is rewarding. A bug in software is very similar. It is a behavior of the software which is not intended. It is often very hard to find bugs in software and can sometimes be attributed to one single character that was wrongly placed. The act of finding bugs is called debugging. Like pest control, the code is searched and cleaned from every bug that influences the behavior negatively. Unfortunately, while fixing software bugs, the programmer accidentally adds other bugs. The Cycle is endless.



10th Fibonacci Number: $13 + 21 = \underline{\underline{34}}$

Loading...



Sometimes a problem can be solved by solving the same problem, but smaller.

And smaller.
And smaller.
Imagine a problem that you can only solve if you know the solution to smaller instances of the same problem. Computing, for example, the 10th element from the famous *Fibonacci sequence*, will require you to sum up the 9th and 8th Fibonacci number. But if you don't know the 10th number, you will most certainly not know the 9th and 8th either. So you'll have to go deeper: In order to obtain the 9th number, you need to add the 8th and 7th element. Knowing the 8th will require the 7th and 6th. And so on. Doing this is called *Recursion*.

But when to stop? Recursion only works if there are certain smallest problems for which you *know* the solution, otherwise you'll keep fragmenting forever. Continuing with the approach above will eventually lead you to computing the 1st and 2nd Fibonacci number, which are defined to be 0 and 1, respectively. Knowing these will allow you to hand the results all the way to the top and finally obtain the desired 10th Fibonacci number.

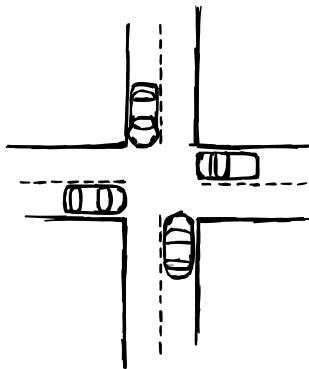
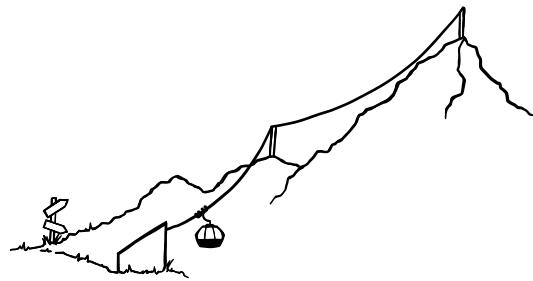
029

Does it scale?

Let's say you want to watch a movie. Downloading it would be a matter of minutes. When ordering it online, the shipment might take a week to arrive. Imagine now you want to watch twenty movies. Downloading them might take a few hours now. But the shipment would still arrive in one week. Now what about a hundred movies? Or a thousand?

When analysing the complexity of an algorithm, it is important to think about the *scalability* of a program. It tells you how the runtime of a program depends on the number of elements processed. In the initial example, the time needed to download the movies depends on the amount. Computer scientists would say: The runtime scales with n , where n stands for the number of movies. When having them shipped, the waiting time doesn't depend on this number, which means it doesn't scale.

030



Being greedy won't make you rich.

031

A very intuitive approach lies at the bottom of so called *greedy algorithms*: At every step of the problem, do what looks best *at that moment*. To imagine this, think about how you hand out change. First, you pick the most valuable coin that covers most of the amount. Then you pick the second largest valued coin of the remaining change. And so on. You'll most definitely be done faster than if you just started assembling random coins.

Greedy algorithms unfortunately only work for rather simple problems. Often it's necessary to plan ahead more than just one step. When climbing a mountain for example, taking at every intersection the path with the steepest slope is a valid approach that will probably bring you to the top. However, by being greedy you most likely missed other opportunities that would have brought you there much faster.

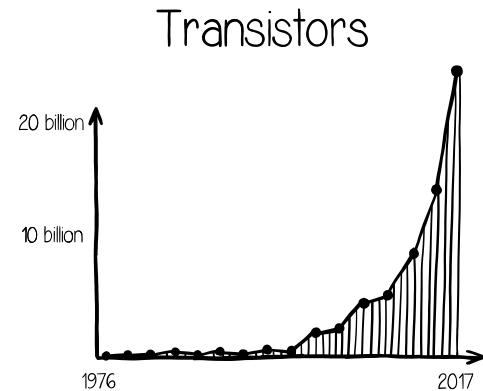
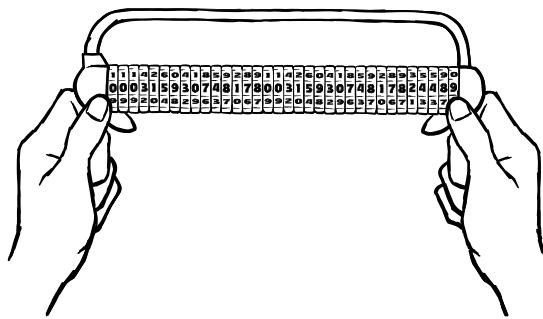
Deadlock

032

Imagine an uncontrolled intersection without any traffic signs. Here, drivers have to respect the priority to the right system. In the situation shown in the left picture, none of the cars can continue. In practice, this is solved by one driver giving up his priority using a hand-sign.

A similar situation can occur when several computer processes running in parallel want to access shared resources. In a *circular wait* scenario, processes block each other from continuing execution as they wait for a resource held by one of their peers.

Since computer processes cannot hand wave at each other, clear rules are needed to avoid this situation in the first place.



"Brute Force" means you are weak

For some problems in Computer Science, it can be proven that there are no "efficient" ways of solving them. In this case, the only resort is **brute force**. Brute force means that you let the computer try out every possible –and we really mean every single, possible and potential– value that solves the problem, until you succeed.

Several cryptographic algorithms are based on such problems where the secret key can only be guessed or found by exhaustively trying out all possible values during a few billion years. Confronted with such time scales, attackers find themselves in a very weak position.

033

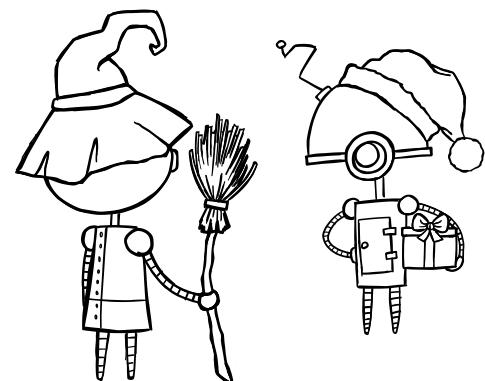
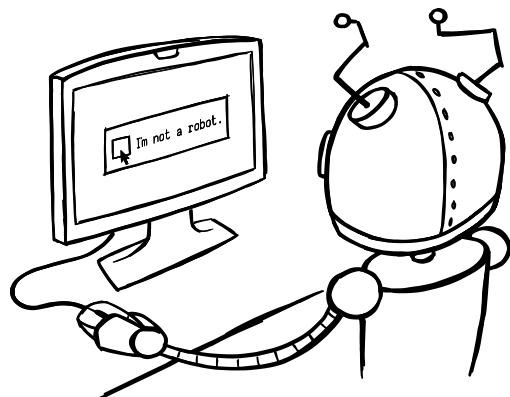
Moore's Law

In 1965, the technological progress lead Gordon E. Moore to formulate his rule of thumb that became famous under the name "Moore's Law". It states that the production costs of microchips used in electronic circuits would fall, leading to smaller and more powerful microchips at a lower overall cost. He estimated, the number of components on a single chip to double roughly every 12 months.

Initially the increase in the number of transistors used in a single chip was responsible for this exponential growth. Later, the downscaling of the components themselves became the driving force.

However, it is unlikely that technological progress will follow Moore's Law forever. As components get smaller and microchips get denser, the manufacturing process will inevitably hit physical limits at some point.

036



Can a computer talk like a human?

Are you sure it is a human answering your question in the chat support of your favorite online store? If not, how you find out? What sounds like an easy question turns out to be rather difficult.

In 1950, a long time before the first online stores opened, the computer science pioneer Alan Turing already thought about this problem and presented the Turing test. The test is passed if a human judge interviewing a human and a machine is unable to distinguish between them based on their answers. The first program that managed to convince some of the judges was ELIZA. It managed to mimic a psychotherapist by reflecting the questions back to the questioner. While there have been several promising attempts in the recent years, yet no program managed to fully pass the test.

However, the next time someone in the chat support of your telecom company tries to calm you down—ask yourself: Am I talking with a human?

038

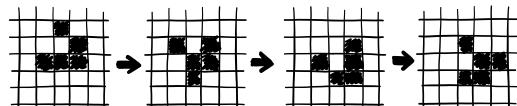
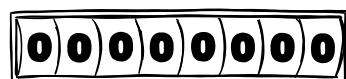
Why does the computer scientist confuse halloween and christmas? Because $25_{dec} = 31_{oct}$

A computer stores everything with 1s and 0s. To encode any number as a string of 1s and 0s you use the binary system. It is similar to the decimal system normal humans use but it has only two digits 1 and 0.

In decimal you have the base 10 and can rewrite the number 25_{dec} as $2 * 10^1 + 5 * 10^0$. In the octal system the number 31_{oct} can be represented as $3 * 8^1 + 1 * 8^0 = 25_{dec}$.

There are also other systems like hexadecimal with the 16 digits 0,1,2,3,4,5,6,7,8,9,A,B,C,E,F and the sexagesimal system with 60 different digits. To get from one system to the other you divide by the base with remainder until you reach 0. Then you string the remainders of these calculations together starting with the one computed last.

039



Overflow

Imagine an odometer of a car. There are six spaces each of which can display a one digit number. Therefore, the highest number that can be represented is 999'999, after which the odometer turns to 000'000 again.

The same thing can happen in a computer. Every number has a fixed amount of bits it can use to store its value. As soon as all of these bits are set, increasing the number again will cause an overflow, meaning that the number will wrap around to its smallest value.

This behavior has led to notorious bugs, such as the software of the Ariane 5 rocket malfunctioning due to an unexpected overflow. This caused a crash shortly after takeoff, costing several hundred million dollars.

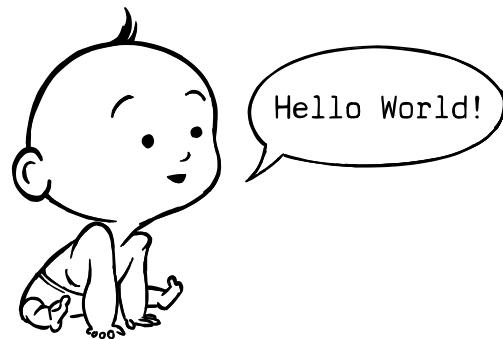
046

Game of life

John Horton Conway invented a simple game with a small amount of rules. Each cell on a square grid is either dead or alive. For each iteration the cell looks at its eight neighbors. If it has less than two living neighbors it dies of loneliness, but if it has more than three neighbors it dies of overpopulation. A cell turns alive if it has exactly three living neighbors.

Those rules are quite simple but the possibilities of interesting simulations are enormous. There is a whole zoo of creatures and self replicating patterns that can evolve in this two dimensional world.

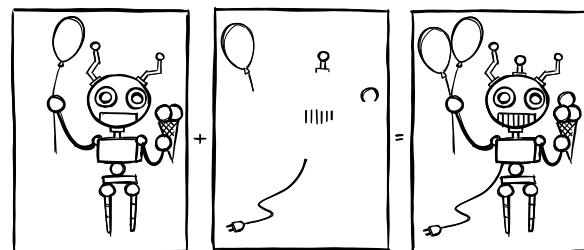
047



The first words of a computer scientist

The first piece of code you usually find in a programming book or tutorial is a "Hello World!" program. The purpose of this program is to display the text "Hello World!". Because this is rather simple for nearly all programming languages, it is a good example to see the basic flavor of a programming language. Due to this tradition, these are usually the first lines of code a programmer writes in their live.

050



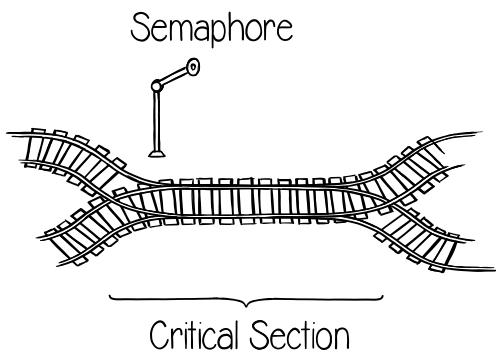
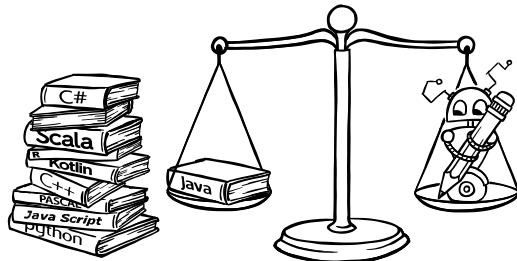
Why Patching is like Picture Puzzles

If you have a buggy version of a program and a corrected one: why send the whole program instead of just the small differences?

This is how updates for your smartphone work: you will receive the set of differences –a so called *patch*– that your smartphone applies to the old and buggy program.

Now go and find the differences: Did we spot them all?

052



Is programming language X more powerful than programming language Y?

There are hundreds or even thousands of programming languages, but are all of these equally powerful, i.e. can they be used to compute an arbitrary problem? The answer is yes — if the language is Turing complete.

A programming language is called Turing complete if it can be used to simulate a Turing machine. Because the Turing machine is able to compute everything that is computable, a Turing complete language is therefore equally powerful.

Most of the commonly used programming languages are Turing complete. But there are some well-known domain specific languages that are not. However, Turing completeness says nothing regarding the performance or the ease of use.

055

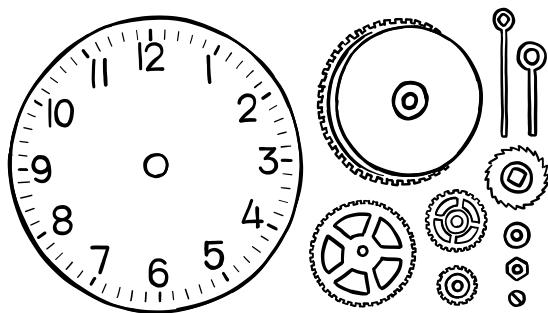
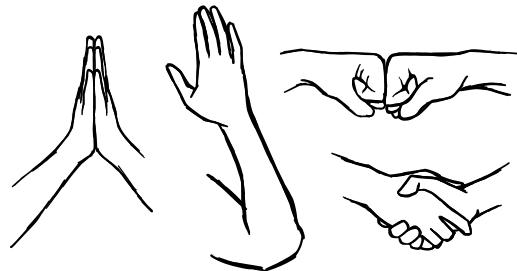
Semaphores prevent crashes

Imagine a railway intersection where only one train can pass at a time. If another train tries to pass the intersection while it's already being used, the trains will crash.

The same thing can happen in a computer. As soon as a resource gets changed from two different sources at the same time, things can go wrong.

In rail traffic, this is solved, by having a *semaphore* to only ever allow one train to pass the intersection. The same is done in computers. Here, a *semaphore* only ever allows one source to access a critical section. Everyone else has to wait, thus preventing crashes or any unintended behavior.

058



Introductions are universal

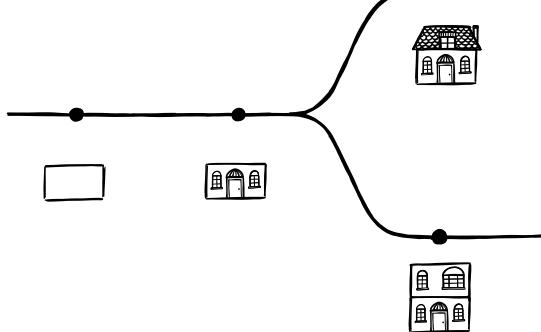
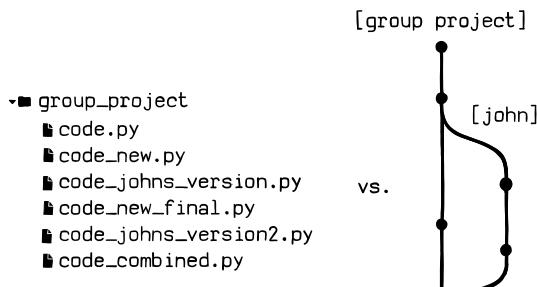
In every society there is some sort of formal introduction when people meet each other for the first time. This is such a fundamental process that it even finds application in computer science, where we call this the handshake. This is the time, where the different processes agree to communicate with each other and clearly state their protocol. Sometimes the handshake is also used to identify each other by means of private keys or passwords.

060

Reverse Engineering

Reverse Engineering describes the process of analyzing an existing apparatus to figure out how it works and how it was made. The same can be done for a computer program, but instead of looking at screws and cogwheels, we are looking at the machine instructions. These are effectively the 0's and 1's that tell the computer what to do. This is obviously much harder to understand than the code that was used to generate the program. A skilled reverse engineer can still use this information to gain knowledge about the software. This can be used to detect security flaws, which can then either be reported to the programmers for them to eliminate or it can be exploited maliciously. Additionally, *Reverse Engineering* is used to change old programs to make them compatible with new hardware or to decipher old file formats.

063



Always use a Version Control System

It's no secret that backing up your data is a good idea, and the same goes for your code. However, keeping a lot of slightly renamed versions of your files for every little modification is probably not the way to go.

Version Control Systems (VCS) allow not only to back up your files, but more importantly help you to keep track of the different versions and what changed in each of them. Even better, multiple people can simultaneously work on one file and then easily combine the different changes.

076

Why do programmers need forks?

Software development is a continuous process. Programmers, engineers, and users, work together to improve the current version of a software.

In open-source projects, the source code of the software is openly available. Everyone can look at it and submit their own contributions. Usually, software projects are happy to receive feedback from the community.

Sometimes, the suggested changes deviate too much from the goal of the project. Developers can then decide to copy the source code and start a new independent project. This deviation from the original project is also called a "fork".

079



??

Being lazy is a good thing

A lazy person is always searching for ways to make things faster and more efficient. Instead of doing the same thing over and over again, a good computer scientist tries to automate the task. This not only improves the reliability, but also saves time in the long run. Furthermore, a good computer scientist is also lazy when it comes to writing code. Rather than solving a specific problem, a good programmer develops a generic solution that can be reused for similar problems.

Good programmers are diligent in order to be lazy. Or with other words: They are strategically lazy. And that's a good thing.

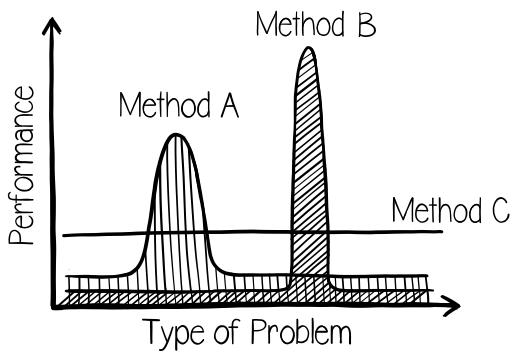
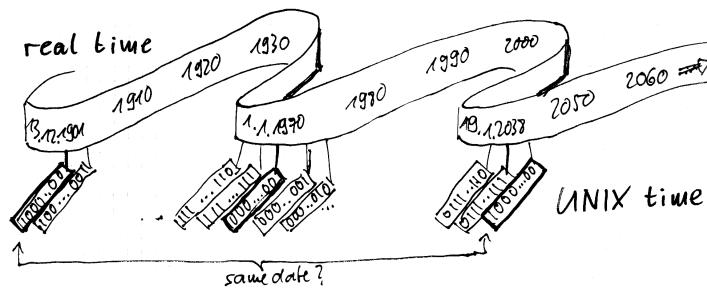
083

Why no Nobel prize for computer scientists?

Some people say it is because that Alfred Nobel's fiancée had an affair with a mathematician, others say because he was more interested in Physics, Chemistry, Medicine, Literature, and Peace than in Mathematics, and computer science is a part of the Mathematics field.

Luckily, there is a "Nobel Prize of Computing" which is the A.M. Turing Award. It is an annual prize given by the Association for Computing Machinery (ACM) to an individual who makes (a) major contribution(s) to the computer field. The award is named after Alan Mathison Turing, who was a mathematician and computer scientist. The first award was at 1966 and went to Alan J. Perlis for his achievements in advanced computer programming techniques and compiler construction. And, the last one was at 2018, and went to Yoshua Bengio, Geoffrey Hinton, and Yann LeCun together, for their improvements in the deep neural networks.

084



No Mercy: The Year 2038 Problem

The date of Jan 1, 1970 is to UNIX developers what year 0 is for Christians, just that in UNIX time you count in seconds, not in days as we do with calendars.

Now here is a problem:

- in UNIX time, two billion seconds after Jan 1, 1970 is Jan 19, 2038
- in UNIX time, two billion seconds before Jan 1, 1970 is Dec 13, 1901

Many UNIX systems can hold time values only up to four billions, and then values start to "wrap around". This means that on Jan 19, 2038, these UNIX systems will think it is the year 1901!

The clock is now ticking for finding all these old UNIX systems and to update their software, giving them more bits for storing *UNIX* time. For sure, some systems will be missed and some systems will be too old to be fixed.

Real time will have no mercy ...

086

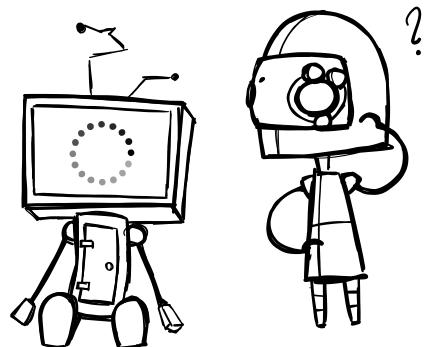
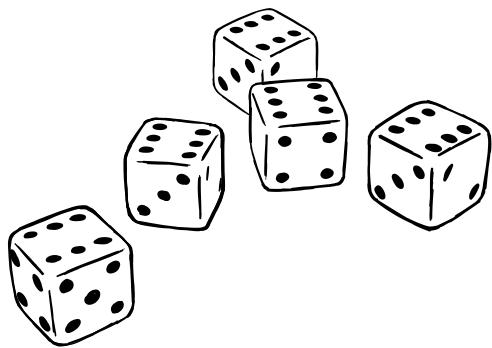
There is no such thing as a free lunch

In Computer Science we have a wide range of very different problems and accordingly many approaches to tackle them. We might for example try to find some 'optimal' solution for a search problem, but it is unclear which specific algorithm one should use. Hence, it is only natural to think about algorithms that could always provide us the best solution, regardless of the problems subtleties.

However, even though a method might result in higher performance for one problem, this comes with the price of diminishing performance for other problems. This is formalized in the so-called 'No Free Lunch' theorems, which state that no single method works better than any other for all possible problems. In practice, it is therefore necessary to carefully select a method for the individual problem at hand.

In other words, there is always a cost associated with selecting a method and unfortunately there is no such thing as a free lunch.

088



The bayesian way of thinking

One might argue that our standard way of reasoning consists of two basic elements: the initial belief we have about a topic and the step of updating this belief when we observe something new.

For example, at first, you would expect a set of dice to be fair when playing a game of Yahtzee. But after observing multiple five-of-a-kind in a row, it is natural to believe that the dice are loaded. And if you actually expected the other players to be sketchy beforehand, you would come to that conclusion way faster.

This simple idea is formalized in Bayes' Theorem, forming the foundation of many Statistics and Machine Learning methods. As a result, it is possible to include expert knowledge into a mathematical model. And even if the used expert knowledge is not totally correct, the model will still be able to overcome that - as long as enough data is observed.

091

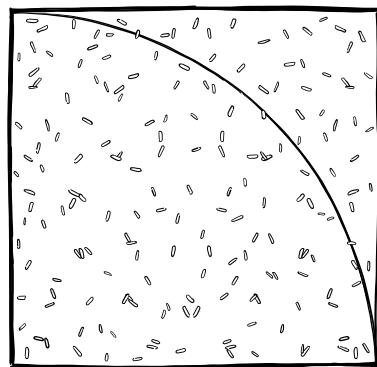
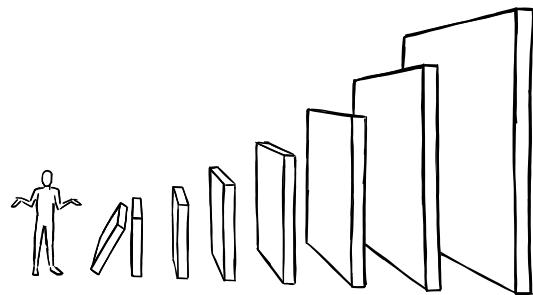
Why There Is No Perfect Anti-Virus

In computer science, simple "Yes/No" questions are called **decision problems**. A decision problem is called **decidable** if we can write a program that can always find an answer for such a problem.

An example for such a decision problem is "does this program contain a computer virus". If this problem was decidable, there could be another program that could always correctly answer this question with "Yes" or "No". A mathematician called Henry Gordon Rice showed that such decision problems about the behavior of another program are always undecidable, unless the answer is **always** the same.

As most viruses are hidden inside other software and do not always expose their real nature, this decision problem can't be solved according to Rice's theorem.

093



Errors propagate!

In daily life it is common to neglect small inaccuracies, be it rounding up change or not being exactly on time.

However, the world of scientific computing is less forgiving. Due to the representation of (floating-point) numbers, there nearly always is a little bit of inaccuracy present and careless computations can increase it substantially. Depending on the problem the errors can increase explosively for each step of an algorithm leading to unanticipated results.

100

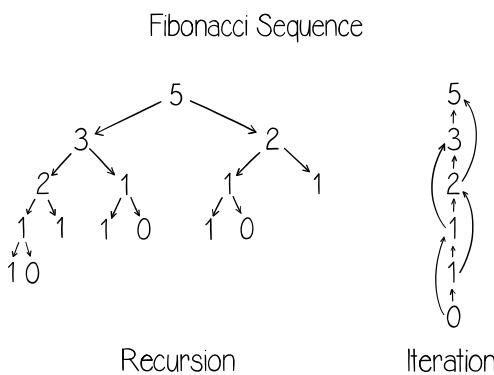
How to calculate Pi with rice

Draw a square and a quatercircle in it on a paper and throw a hand full of rice on it. Done!

Well almost done. Now you count how many grains of rice landed inside the square and how many landed inside the quatercircle. The ratio of these two numbers is an approximation of the ratio of the two areas $\pi r^2 / 2 = \pi/4$. Multiply the ratio by 4 and you get Pi, well an approximation. If you are not happy with the accuracy just use more rice.

This is one example for the Monte Carlo method. Key is to avoid computing a complicated formula by using random samples in a clever way.

104



Oops! My computer keeps freezing, could it be a livelock?

Yes, It could be. A livelock is one of the reasons that might freeze your system. But, what is a livelock?

Suppose two gentlemen are sitting in a restaurant and the waiter brings the order without enough cutlery. The first gentleman has a fork and the second one has a knife. To be able to eat, the first asks the second for the knife and at the same time the second asks the first for the fork. They exchange the cutlery, and the situation happens again; the first has the knife and the second has the fork. They keep exchanging the cutlery without eating.

This situation could happen in operating systems as well. When process A holds resource R1 and requires resource R2 which is held by process B and vice versa. They request the required resource from each other. Then, they exchange the resources and keep exchanging them without completing their tasks.

105

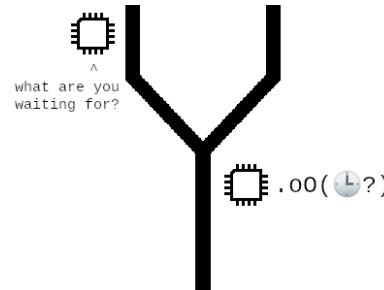
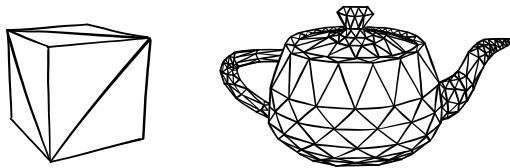
Why solve any problem with recursion when iteration works just as well?

Solving a problem step by step, starting at the bottom and linearly working towards the solution, is called iteration. Recursion will do it the other way around. It will start at the top, search its path down and then hand the obtained results all the way back. Sounds tedious, right? And since every problem solved by recursion could also be solved by iteration way faster, why use recursion at all?

The beauty and usefulness of recursion lies in its intuitive nature. When using iteration, you need to code every step towards the solution yourself. When using recursion, you only need to think about the pattern of the problem. The computer will automatically apply this pattern, again and again, until it reaches the bottom of all the smaller parts of the problem. Then it will combine everything to obtain the solution.

There is no general rule when to use recursion and when not to. For any given problem, simply pick the approach that seems the most natural to you.

108



World of triangles

The 3D pictures seen in computer animated films or computer games consist of many small triangles. The shape of a figure in such an image comprises many points. To visualize a surface, a set of those points is connected by lines. The mathematical figure which emerges through this process is called polygon. The smallest polygon possible is the triangle, where only three points are connected. Furthermore, a triangle guarantees that all the points of the polygon are on the same plane.

The more triangles are used, the more detailed a surface is. Whereas a simple object with flat surfaces like a cube can be modeled by using only a few triangles, a more complex structure like a teapot requires many more triangles for representing its round shape.

109

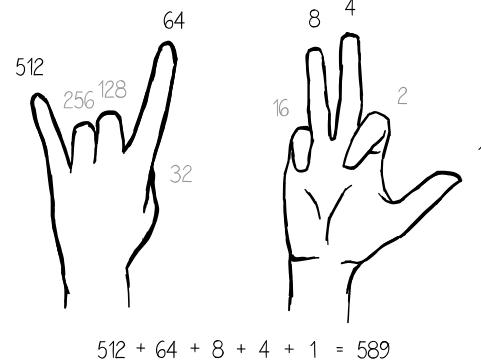
Guess Fast Or Wait!

Imagine you're in an unfamiliar city and trying to find a new place. Also, you're using your smartphone for navigation. And because you have a good sense of direction, you already start walking the way you think is correct while looking up the way on your phone.

If you guessed the direction correctly, you didn't waste any time waiting for the result, and can continue on your way. If your guess was wrong, you have to turn around.

A modern CPU does this as well: When its course of action depends on a result it would have to wait for, it makes an educated guess and continues along this path. This is called **speculative execution**. If the guess was wrong, the CPU has to discard all the actions it has done in error. It did not waste any time though, as it would have spent the time waiting instead. However, if the guess turns out to be correct, the CPU just gained a little speed boost.

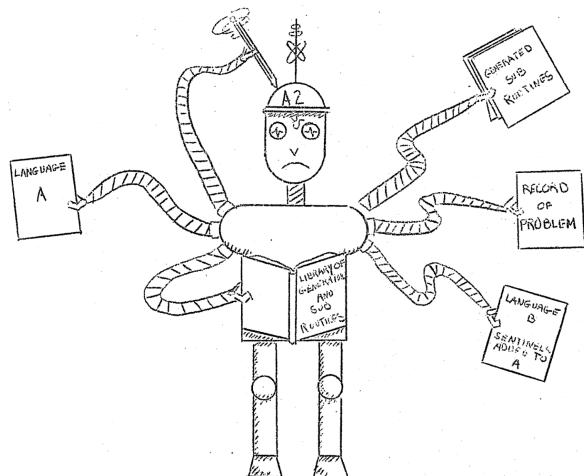
110



Counting to more than ten with two hands

Normally you count with your fingers by stretching them out one after another. You don't care about which ones are stretched out, but only about the amount of them. Each of them is as significant as the other. By changing the significance in the way that the right thumb has significance 1 and the left neighbor neighbor is always twice as significant we can get 1024 different numbers. $1+2+4+8+16+32+64+128+256+512=1023$ and 0. This is called binary counting and exactly the way a computer counts but with zeros and ones instead of bent and stretched fingers.

111



We have Machines within Machines

In 1954, computers just computed numbers. Grace Hopper, a famous female programmer around then, had a hard time to convince her colleagues that these machines can do more.

The picture is hers and shows how a computer program can translate one computer program into another one. She even coined the term that is still in use today: a **compiler**.

Nowadays, we have many more of these translation and transformation machines, with names as funny as the picture: *cross-compilers*, *transpilers*, and *compiler-compilers* (*cocos*). Their mission is to translate from one computer language to another, for example to "compile" a human-written program to machine code, or to "transpile" from a new version of a programming language to the old one.

112



Code:	Comments:
++++++ [>+++++>++++++>++>+<<<-] >+. >. +++++. . . +++. >+. <<+++++++. >. +++. ----. -----. >+. >. +++. 	start of loop end of loop H e l o space W o r l d ! line feed new line

Why Uptime Matters

The time since a computer was started is referred to as its **uptime**. When the computer is rebooted (or crashes), its uptime is reset to zero. The time the computer is not running is called **downtime**.

When running a **service**, such as a website, we want it to be available at all times. This means that the **server** the website is running on has to be available all the time. As a high uptime means that a server has been running a long time without interruption, it can indicate that the provided services are also available with few interruptions.

114

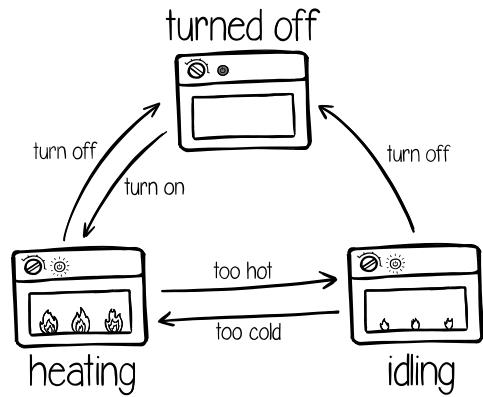
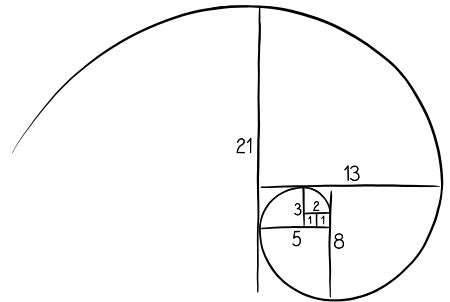
Esoteric Science

There are thousands of programming languages. But some of them are just laborious and inefficient. They are not designed for practical usage but to implement different aspects of unusual language concepts. Those programming languages are called esoteric. One of them is called Brainfuck. It consists of the eight language commands:

< > + - , , []

To be able to understand the code, more comments are needed to explain the code than the size of the code itself. It is possible to write any program with this language. But the question is, if one really wants.

116



Why stick to one example

Using the same example, again and again, seems to be quite boring without any added benefit. In some instances, however, this can be used for some positive effect. Using a known example to show different things (like methods or algorithms) saves time, as the people don't have to think about the example used in much detail. In this book, the common example is the Fibonacci Sequence which starts with 0 and 1. Every next number is the sum of their two predecessors.

0, 1, 1, 2, 3, 5, 8, 13, 21, ...

The well-known picture according to this Sequence is the Fibonacci spiral.

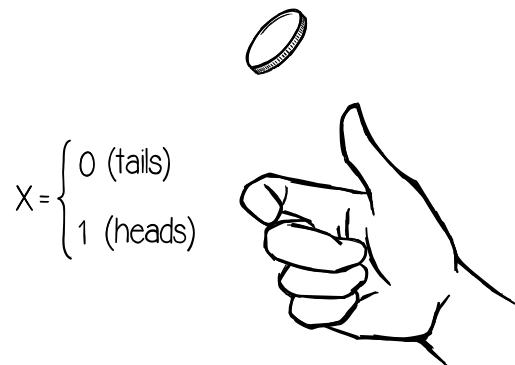
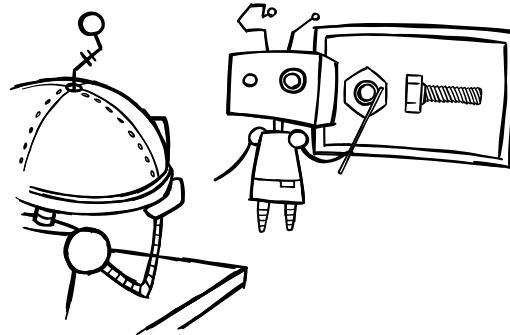
117

Think about the different states of your system

State machines are a concise way of describing systems with clearly distinguishable phases or states. The concept of state machines were developed in the context of computer systems but is now also used by engineers from other disciplines.

A state machine is a system which defines its current status through a set of states. The current state determines what the system is doing. In case of the example, this means whether the heating element is activated or not. The transition between the states happens based on events. These events can be triggered by external factors like a button pressed by the user or a sensor reading but can also be internal like the result of a computation.

118



How to find a solution through learning

If we try to solve a problem, we always appreciate if we get feedback on how good we were and this is the same case for a program. In supervised learning we first let a program solve problems, where we do know the results, like for example finding the labels of different points, so called classification. Through this training we know how good the solutions of the program are and can adjust the program for efficiency, and can then use it for the same kind of problem where we do not know the solution.

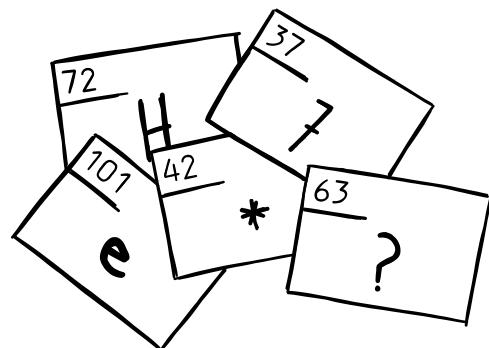
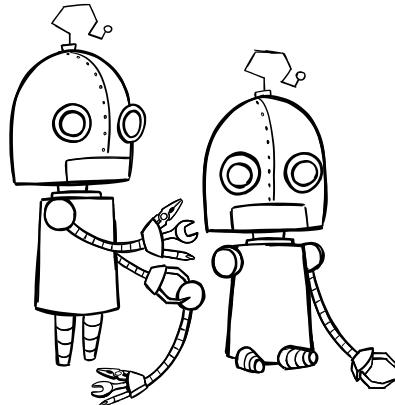
119

Random Variables

Random variables quantify the outcomes of a random process which could have different outputs. They are different from algebraic variables that only have a single value. In our daily life, there are many examples of such random processes, such as the stock market, blood pressure, or even simpler things such as flipping a coin or rolling a dice.

To be able to do some mathematics on these outcomes, we need to define a value for each respective one. For instance, according to our example, we could check the probability of X being head by calculating $P(X=1)$. In probability theory, many methods use random variables.

120



The fixpoint of describing yourself

What could be easier than to write a program that produces its own code as output? It's not that easy — even a versed programmer will spend some hours to figure out the trick (which can be different for each programming language).

One trick is that the program must contain a description of itself, inside its own code. And because outputting that self-description also needs code, this self-description must be somehow compressed.

But once written, such a program (also called a "Quine") reproduces itself each time it is run, and this in turn is called a fixpoint. One could call this the fixpoint of life because any lifeform must master this trick - including robots if they wanted to build themselves and become a lifeform.

121

Every letter is a number

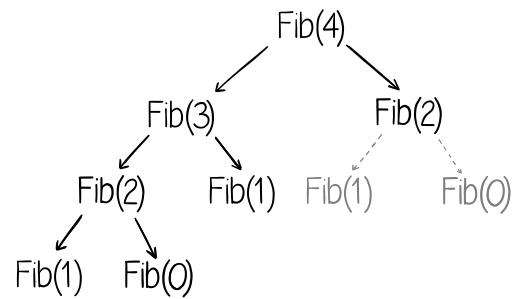
In computers, letters are represented as numbers. To encode them, every letter is assigned its own decimal number between 0 and 127. The "H", for example, is encoded as the 72. The string "Hello world" is stored as a sequence of numbers:

H	e	l	l	o		s	p	a	t	w	o	r	l	d
72	101	108	108	111	32	119	111	114	108	100				

As everything is stored in binary in computers, those numbers are translated in binary code. Not only the capital letters are encoded in this way, but also the lowercase letters, the digits 0-9 and punctuation symbols.

This very common character encoding is called ASCII (American Standard Code for Information Interchange). Today, the international standard for encoding is called Unicode. It is way bigger than ASCII — big enough that every possible character of every language may have its own digital code.

122



Try dynamic programming with recursive problems

Dynamic programming is similar to recursion as it requires to divide the problem into sub-problems. By solving the sub-problems, we get the solution of the bigger one. The difference is that in recursion we repeatedly calculate the results, but in dynamic programming, we memorize the old results, so we do not calculate the results again and again.

In Fibonacci series for instance, while calculating Fib(5), we know that Fib(0)=0 and Fib(1)=1, so we start with them to calculate Fib(2). By memorizing Fib(2), we can calculate Fib(3) which is equal to Fib(2)+Fib(1). Then, memorizing Fib(3) and calculating Fib(4)=Fib(3)+Fib(2). We continue in the same way until we find our target.

123

17_dec



21_oct

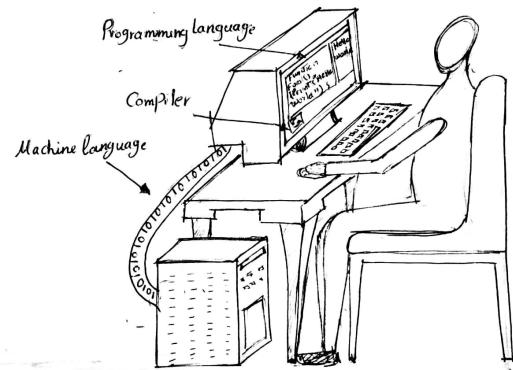


What if humans had 8 fingers instead of 10?

Whenever we are dealing with numbers in our daily lives, we use the so called *decimal system*. This means that our entire counting system is based on the number 10. A given number, for example 17, actually means „1 time the number 10, plus 7 times the number 1“. It's the same as counting with fingers. Once you reached 10, you keep the result in your mind and start over with one finger again.

Now what if we had 8 fingers instead? Our counting system might be based on the number 8. This would be called *octal system*. Our decimal number 17, also called [17_dec] to prevent confusion, would instead be written as [21_oct], which actually stands for “2 times the number 8, plus 1 time the number 1”.

124



Machine Language vs. Programming Language

A programming language is a language that programmers use to develop applications. It is also called high-level language as they are more human readable. Examples of such programming language include -but not limited to- C++, Java, and Fortran.

On the other hand, machine language is a low-level language, because the computer can understand and process directly without any transformations. It is only one language which is a stream of 0's and 1's.